

Class 7: Machine Learning 1

Nicole Lin (PID: A17826971)

Background


Today we will begin our exploration of some important machine learning methods. Namely **clustering** and **dimensionality reduction**.

Let's make up some input data for clustering where we know what the natural "clusters" are.

The function `rnorm()` can be useful here.

```
hist( rnorm(30, mean = 3))
```

Histogram of `rnorm(30, mean = 3)`




Q. Generate 30 random numbers centered at +3 and 30 random numbers centered at -3.

```

tmp <- c(rnorm(30, 3),
         rnorm(30, -3))

x <- cbind(x=tmp, y=rev(tmp))
plot(x)

```



K-means clustering

The main function in “base R” for K-means clustering is called `kmeans()`.

```

km <- kmeans(x, centers = 2)
km

```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.075108	-3.224556
2	-3.224556	3.075108

Clustering vector:

```
Within cluster sum of squares by cluster:  
[1] 53.53667 53.53667  
  (between_SS / total_SS =  91.7 %)
```

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. What component of the results object details the cluster sizes? size

km\$size

[1] 30 30

Q. What component of the results object details the cluster centers?

km\$centers


	x	y
1	3.075108	-3.224556
2	-3.224556	3.075108

What component of the results object details the cluster membership vector (i.e. our main result of which points lie in which cluster)?

km\$cluster

Q. Plot our clustering results with points colored by cluster and also add the cluster centers as new points colored blue.

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15)
```



Q. Run `kmeans()` again and this time produce 4 clusters (and call your result object `k4`) and make a results figure like above.

```
k4 <- kmeans(x, 4)  
k4
```

K-means clustering with 4 clusters of sizes 6, 10, 30, 14

Cluster means:

	x	y
1	-3.694187	2.098295
2	-3.977333	3.931372
3	3.075108	-3.224556
4	-2.485587	2.882126

Clustering vector:

Within cluster sum of squares by cluster:


```
[1] 5.570452 12.290363 53.536665 7.462507
```

(between_SS / total_SS = 93.9 %)

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(x, col=k4$cluster)  
points(k4$centers, col="blue", pch=15)
```



The metric

```
km$tot.withinss
```

```
[1] 107.0733
```

```
k4$tot.withinss
```

```
[1] 78.85999
```

Q. Let's try different number of K (centers) from 1 to 30 and see what the best result is.

```

i <- 1
ans <- NULL
for(i in 1:30) {
  ans <- c(ans, kmeans(x, i)$tot.withinss)
}

ans


```

```

[1] 1297.646423 107.073330 84.796798 78.498319 72.022830 49.597681
[7] 42.091682 36.589927 32.475752 28.362635 28.413756 19.967946
[13] 14.023506 21.418543 14.346024 15.484308 8.669601 9.096797
[19] 10.440591 5.004807 5.178799 7.660381 4.475855 4.745259
[25] 3.912805 3.372393 2.795203 2.993818 2.877597 2.257538

```

```
plot(ans, typ="o")
```




Key point: K-means will impose a clustering structure on your data even if it is not there - it will always give you the answer you asked for even if that answer is silly! The best result is the one with 2 clusters (first one that drops off the “cliff”) because it has the fewest number of clusters but the smallest tots.withinss, meaning that point has clusters that are the tighter and better-defined while still having a small number of clusters.

Hierarchical Clustering

The main function for Hierarchical Clustering is called `hclust()`. Unlike `kmeans()` (which does all the work for you), you can't just pass `hclust()` our raw input data. It needs a "distance matrix" like the one returned from the `dist()` function.


```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

Cluster Dendrogram



To extract our cluster membership vector from a `hclust()` result object we have to "cut" our tree at a given height to yield separate "groups/branches."

```
plot(hc)
abline(h=8, col="red", lty=2)
```



To do this, we use the `cutree()` function on our `hclust()` object:

```
grps <- cutree(hc, h=8)  
grps
```

```
table(grps, km$cluster)
```

grps	1	2
1	30	0
2	0	30

PCA of UK food data

Import the data set of food consumption in the UK:

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

One solution to set the row names is to do it by hand...

```
rownames(x) <- x[,1]
```

To remove the first column, I can use the minus index trick.

```
x <- x[,-1]
x
```

		England	Wales	Scotland	N.Ireland
Cheese		105	103	103	66
Carcass_meat		245	227	242	267
Other_meat		685	803	750	586
Fish		147	160	122	93
Fats_and_oils		193	235	184	209

Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

A better way to do this is to set the row names to the first column with `read.csv()`.

```
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Prefer the second method because it doesn’t take away a column each time you run the code.
Also is shorter.

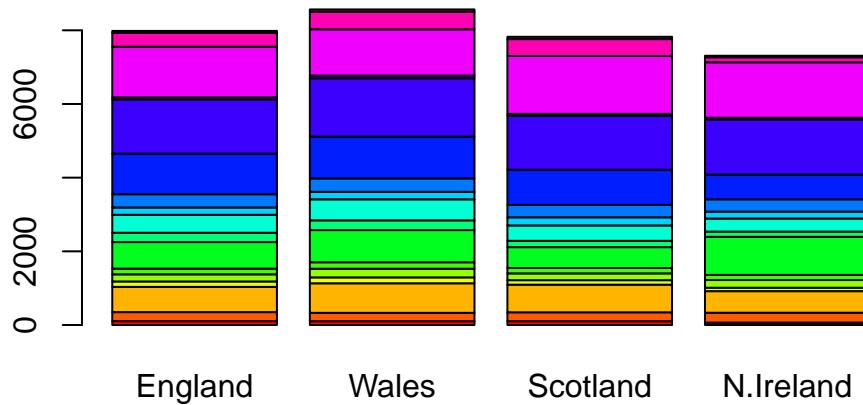
Spotting major differences and trends

It's difficult even in this wee 17 dimention data set...

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```




Q3: Changing what optional argument in the above barplot() function results in the following plot?


Changing the “beside” argument from T to F changes it from a histogram graph to a stacked graph.

Pairs plots and heatmaps

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



PCA to the rescue

The main PCA function in “base R” is called `prcomp()`. This function wants the transpose of our food data as input (i.e. the foods as columns and the countries as rows).

```
pca <- prcomp( t(x) )
```

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.7e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.0e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.0e+00

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

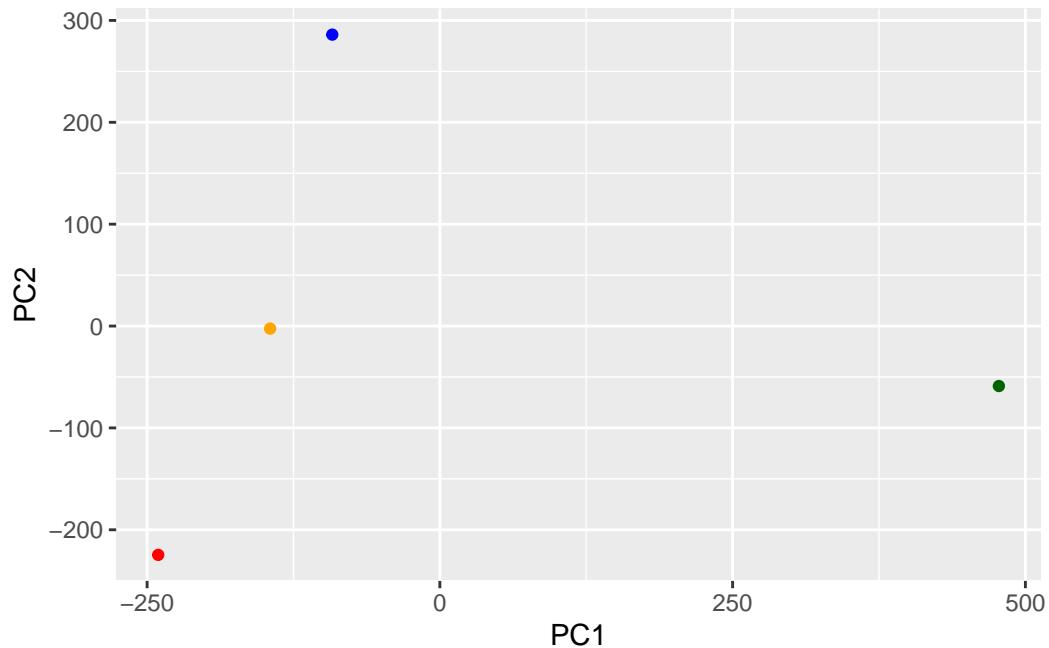
$class
[1] "prcomp"
```

To make one of our main PCA result figures, we turn to `pca$x` the scores along our new PCs. This is called “PC plot” or “score plot” or “ordination plot”...

```
my_cols <- c("orange", "red", "blue", "darkgreen")
```

```
library(ggplot2)

ggplot(pca$x) +
  aes(PC1, PC2) +
  geom_point(col=my_cols)
```



The second major result figure is called a “loadings plot” of “variable contributions plot” or “weight plot”.

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```

